

My recommendation is not to buy AI for this.

Plain rules already assign 60 of 111 vehicles and stop on the rest with a named reason. What I'm asking for are two decisions and a capped test.

TODAY

60 of 111 assigned

No model, no network calls, no licence.
Every decision is inspectable row by row.

YOUR DECISION

30 more are waiting

Two one-time rulings. No budget, no build, and not mine to make.

THE ASK

One capped pilot

Run beside the current process.
Nothing published changes until it passes.

Every vehicle falls into one of four buckets. Only one of them is genuinely unanswerable.

VIO accuracy is the constraint. Manual effort is the cost. This splits them.

NO DECISION NEEDED

60 vehicles

The evidence separates exactly one candidate. A reviewer confirms rather than investigates.

ONE RULING UNLOCKS MANY

30 vehicles

Blocked on two catalogue questions. Decide each once and it applies to every future batch.

ONE DATA CORRECTION

7 vehicles

A registry capacity that contradicts every other field on the same row, including its own VIN.

NO ANSWER EXISTS

14 vehicles


The RDM reference holds no version at the registered capacity and power. It's not a mapping failure; it's a catalogue gap.

Two of the four buckets are answered once by the data owner (the thirty and the seven), and those answers apply to every future batch.

60 vehicles assigned.
51 stopped with a named reason.

54.1%

Distinct vehicles with an accepted kType



Each square is one distinct vehicle ID. 112 worksheet rows, one exact duplicate.

Accepted kType assignment	60
All candidates contradicted	40
Candidates not separated	10
Conflicting source data	1

Accepted: 20 Hyundai Kona → 129021 · 18 Mercedes C-Class 2.0P → 137963
13 Hyundai Tucson → 115220 · 9 Nissan X-Trail Hybrid → 124055

54.1% is coverage, not accuracy. Two supplied labels cannot prove that these 60 assignments are right.

Nothing is silently dropped. Every stop names a field and both values.

A visible exception is a task you can route. Sometimes the registry is wrong, and the evidence says so.

20 NISSAN X-TRAIL

One kilowatt apart

Registry says 125 kW. Catalogue says 126 kW. Everything else agrees.

14 FORD RANGER

No version in the extract

Registry says 1996 cc and 157 kW.
The 365-row extract's 13 Rangers hold 2184–3198 cc, 57–147 kW.

10 HOLDEN COLORADO

Two types, one body word

kTypes 100868 and 100871 are identical on every technical field. Only the body descriptor differs.

Seven Mercedes rows look like a registry error, not a mapping failure: power, fuel, variant and engine agree with one kType; only the recorded capacity disagrees, 1491 cc against 1991.

A looser algorithm would have assigned nearly all 111.

My first attempts did exactly that. That was the problem, not the achievement.

APPROACH	WHAT IT DID	ASSIGNED
First experiments	Nearest candidate; every row gets an answer	Nearly all 111
v1 · first structured pass	Make, model and specification comparison	51 · 45.9%
v2 · grounded the rules	Removed every correspondence no publisher states	42 · 37.8%
v3 · current	Added sourced identifier evidence to expand the VIN decoder profiles	60 · 54.1%

Switch the identifier evidence off and coverage **rises** to 80 of 111. Twenty of those are a previous-generation X-Trail agreeing on capacity, power, fuel and drive, and contradicted by the VIN. Twenty confident errors.

Thirty vehicles are not blocked by code. They are blocked by two unanswered questions.

Neither is mine to decide. Both are one-time answers that apply to every future batch.

Ruling 1 · +20 vehicles

Does a 1 kW gap count as agreement?

You have seen the gap. The question is whether two publishers one kilowatt apart describe the same vehicle. I verified that allowing it leaves exactly one compatible candidate per row, but I will not let an inferred tolerance manufacture evidence. The ruling is yours.

If ruled: 20 X-Trail resolve. 80 of 111, or 72.1%.

Lookup 2 · +10 vehicles

What does TecDoc mean by "Platform/Chassis"?

Not a judgement call, but a definition TecAlliance already holds. NZTA answers its half: it records these as Utility and keeps Cab/chassis as a separate code. The other half is a lookup, not a decision.

If ruled: 10 Colorado resolve. 90 of 111, or 81.1%.

This is built, not proposed.

A reproducible command, a reviewer interface over the same policy, and an audit trail for every row.

PIPELINE

One command

Reads both workbooks, writes the result and a run report. No network calls, no model at runtime.

REVIEW

Local reviewer interface

Upload a registry file, inspect totals, open candidate-level evidence, download the workbook.

AUDIT

Nine-sheet workbook

Results, candidate evidence, review queue, duplicates, VIO counts, active policy, input hashes.

96 tests and 69 subtests pass, and the prototype was run against the published 2020 NZ register (181,790 vehicles, sourced separately, because the supplied extract is 111 rows) for stability and throughput. Neither establishes mapping accuracy.

The next investment buys a decision.

Run it beside the current process. One representative batch, independent adjudication, nothing published changes.

UNPROVEN

Are the 60 correct?

Two supplied labels agree and cover one kType. That cannot estimate precision on anything.

UNPROVEN

Does total effort fall?

Coverage is not labour saved. Checking and rework have never been timed against today's process.

What I need: a named business owner, an end date, a spending cap and about 150 expert decisions. Identical vehicles only have to be judged once, so 150 judgements cover what would otherwise be 1,400 reviews.

Agree what success means before we see the results.

Quality is a condition of expansion. Savings must survive the cost of checking.

QUALITY

A precision floor, set in advance

Measured on adjudicated **decisions**, not rows, and counting valid matches missed, not only errors among the accepted.

PEOPLE

Reviewer time, measured both ways

Review and rework timed against the current process, including the unresolved queue, not coverage as a proxy.

OPERATION

Throughput, already largely answered

That same published register, 181,790 vehicles, mapped in 36 seconds. The pilot confirms it at your volumes and on your infrastructure.

All three pass, we expand. One fails and is fixable, we revise within the cap. Otherwise we stop and keep the findings.

The current process stays in place until the alternative earns its role.

Three practical options. I am recommending the middle one.

CONTINUE TODAY

Lowest immediate change

Defensible if today's effort and turnaround already meet the need. But it answers nothing new, and whatever the review of these vehicles costs you today, you keep paying it, including on the 60 the evidence already separates.

RECOMMENDED

Bounded parallel pilot

The only option that converts an opinion into a measurement. You set the ceiling, and the answer arrives with the evidence behind it.

ROLL OUT NOW

Earlier operational change

Premature. Accuracy and savings have not been established on any representative batch.

Where AI would have to earn its place.

I opened by recommending against AI here. This is what would change my mind, and what it would have to prove first.

NOT YET

A trained matcher

Becomes possible once adjudicated labels exist. Today there are two, so there is nothing to train or calibrate on.

SEPARATE TEST

Reviewer assistance

A language model could explain difficult records. Keep it only if checked answers measurably reduce total effort.

UNCHANGED

The same quality gate

Either addition faces the identical precision floor. Neither gets a lower bar for being newer.

Two rulings this week.

One capped test before we commit to anything.

Neither requires budget approval. Both produce evidence you can act on.

THIS WEEK · FREE

Name the data owner

Who rules on the 1 kW gap, and who defines "Platform/Chassis"? Thirty vehicles follow.

THEN · CAPPED

Sponsor the pilot

The terms are already specified. What I need from you is the sponsorship and the reserved review time. The cap means the downside is known before it starts.

ONCE · DECISIVE

One evidence review

One meeting at the end, not rolling updates. You decide then: expand, revise or stop. The findings are yours whichever you choose.